

# Deep Learning: Homework 2

Instructed by *Yi Wu*

Due on Mar 24, 2021

**Runlong Zhou** YaoClass 82 2018011309

## Model Details

The core part of my model is adapted from VGG16. My model takes a 3-channel  $96 \times 96$  picture as input and outputs a 10-dim vector as classification result.

The model consists 5 parts, transforming feature sizes as:  $96 \times 96 \times 3 \rightarrow 48 \times 48 \times 64 \rightarrow 24 \times 24 \times 128 \rightarrow 12 \times 12 \times 256 \rightarrow 6 \times 6 \times 512 \rightarrow 3 \times 3 \times 512$ . In each part, there are two to three groups of (conv2d, batchnorm2d, relu), followed by a final maxpool2d whose stride is 2 to make side length halved.

The final output is calculated by a full-connection with  $3 \times 3 \times 512$  inputs and 10 outputs.


## Hyper-Parameters

1. num-epoch: default 100.
2. lr: default 5e-3.
3. weight-decay: default 0.
4. batch-size: default 16.
5. aug-prob: probability to do augmentation, default 0.5.
6. ensemble: usage specified in the tricks, default=6.

## Tricks

1. **Data augmentation:** Pepper-noise (deprecated), Scale-then-randomcrop (substituted by ensemble), Random-horizontal-flip, Random-vertical-flip (deprecated), Random-rotation (deprecated).
2. **Ensemble:** Use several (number specified by `-ensemble`)  $96 \times 96$  crops of the original  $128 \times 128$  picture as input, and average softmax(outputs) as ensemble probability. To create the same loss function, first apply `log()` then use `NLLLoss()`.

## Training Curve



Best model is at epoch 88.

## Modifications to evaluation.py

1. Line 5: Net  $\rightarrow$  \*.
2. Line 35: map\_location=device, since my model is not trained on CPU or GPU0.